

Data Manipulation Challenge

A Mental Model for Method Chaining in Pandas

Data Manipulation Challenge - A Mental Model for Method Chaining in Pandas

! Challenge Requirements In Section [Student Analysis Section](#)

- Complete all discussion questions for the seven mental models (plus some extra requirements for higher grades)

! Note on Python Usagew

Recommended Workflow: Use Your Existing Virtual Environment If you completed the Tech Setup Challenge Part 2, you already have a virtual environment set up! Here's how to use it for this new challenge:

1. **Clone this new challenge repository** (see Getting Started section below)
2. **Open the cloned repository in Cursor**
3. **Set this project to use your existing Python interpreter:**
 - Press **Ctrl+Shift+P** → “Python: Select Interpreter”
 - Navigate to and choose the interpreter from your existing virtual environment (e.g., `your-previous-project/venv/Scripts/python.exe`)
4. **Activate the environment in your terminal:**
 - Open terminal in Cursor (‘Ctrl + “)
 - Navigate to your previous project folder where you have the `venv` folder
 - **Pro tip:** You can quickly navigate by typing `cd` followed by dragging the folder from your file explorer into the terminal
 - Activate using the appropriate command for your system:
 - **Windows Command Prompt:** `venv\Scripts\activate`

- **Windows PowerShell:** `.\venv\Scripts\Activate.ps1`
 - **Mac/Linux:** `source venv/bin/activate`
 - You should see (venv) at the beginning of your terminal prompt
5. **Install additional packages if needed:** `pip install pandas numpy matplotlib seaborn`

Cloud Storage Warning

Avoid using Google Drive, OneDrive, or other cloud storage for Python projects! These services can cause issues with: - Package installations failing due to file locking - Virtual environment corruption - Slow performance during pip operations

Best practice: Keep your Python projects in a local folder like `C:\Users\YourName\Documents\` or `~/Documents/` instead of cloud-synced folders.

Alternative: Create a New Virtual Environment If you prefer a fresh environment, follow the Quarto documentation: <https://quarto.org/docs/projects/virtual-environments.html>. Be sure to follow the instructions to activate the environment, set it up as your default Python interpreter for the project, and install the necessary packages (e.g. pandas) for this challenge. For installing the packages, you can use the `pip install -r requirements.txt` command since you already have the requirements.txt file in your project. Some steps do take a bit of time, so be patient.

Why This Works: Virtual environments are portable - you can use the same environment across multiple projects, and Cursor automatically activates it when you select the interpreter!

The Problem: Mastering Data Manipulation Through Method Chaining

Core Question: How can we efficiently manipulate datasets using `pandas` method chaining to answer complex business questions?

The Challenge: Real-world data analysis requires combining multiple data manipulation techniques in sequence. Rather than creating intermediate variables at each step, method chaining allows us to write clean, readable code that flows logically from one operation to the next.

Our Approach: We'll work with ZappTech's shipment data to answer critical business questions about service levels and cross-category orders, using the seven mental models of data manipulation through `pandas` method chaining.

AI Partnership Required

This challenge pushes boundaries intentionally. You'll tackle problems that normally require weeks of study, but with Cursor AI as your partner (and your brain keeping it honest), you can accomplish more than you thought possible.

The new reality: The four stages of competence are Ignorance → Awareness → Learning → Mastery. AI lets us produce Mastery-level work while operating primarily in the Awareness stage. I focus on awareness training, you leverage AI for execution, and together we create outputs that used to require years of dedicated study.

The Seven Mental Models of Data Manipulation

The seven most important ways we manipulate datasets are:

1. **Assign:** Add new variables with calculations and transformations
2. **Subset:** Filter data based on conditions or select specific columns
3. **Drop:** Remove unwanted variables or observations
4. **Sort:** Arrange data by values or indices
5. **Aggregate:** Summarize data using functions like mean, sum, count
6. **Merge:** Combine information from multiple datasets
7. **Split-Apply-Combine:** Group data and apply functions within groups

Data and Business Context

We analyze ZappTech's shipment data, which contains information about product deliveries across multiple categories. This dataset is ideal for our analysis because:

- **Real Business Questions:** CEO wants to understand service levels and cross-category shopping patterns
- **Multiple Data Sources:** Requires merging shipment data with product category information
- **Complex Relationships:** Service levels may vary by product category, and customers may order across categories
- **Method Chaining Practice:** Perfect for demonstrating all seven mental models in sequence

Data Loading and Initial Exploration

Let's start by loading the ZappTech shipment data and understanding what we're working with.

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta

# Load the shipment data
shipments_df = pd.read_csv(
    "https://raw.githubusercontent.com/flyaflya/persuasive/main/shipments.csv",
    parse_dates=['plannedShipDate', 'actualShipDate']
)

# Load product line data
product_line_df = pd.read_csv(
    "https://raw.githubusercontent.com/flyaflya/persuasive/main/productLine.csv"
)

# Reduce dataset size for faster processing (4,000 rows instead of 96,805 rows)
shipments_df = shipments_df.head(4000)

print("Shipments data shape:", shipments_df.shape)
print("\nShipments data columns:", shipments_df.columns.tolist())
print("\nFirst few rows of shipments data:")
print(shipments_df.head(10))

print("\n" + "="*50)
print("Product line data shape:", product_line_df.shape)
print("\nProduct line data columns:", product_line_df.columns.tolist())
print("\nFirst few rows of product line data:")
print(product_line_df.head(10))

```

Shipments data shape: (4000, 5)

Shipments data columns: ['shipID', 'plannedShipDate', 'actualShipDate', 'partID', 'quantity']

First few rows of shipments data:

	shipID	plannedShipDate	actualShipDate	partID	quantity
0	10001	2013-11-06	2013-10-04	part92b16c5	6
1	10002	2013-10-15	2013-10-04	part66983b	2
2	10003	2013-10-25	2013-10-07	part8e36f25	1
3	10004	2013-10-14	2013-10-08	part30f5de0	1
4	10005	2013-10-14	2013-10-08	part9d64d35	6

5	10006	2013-10-14	2013-10-08	part6cd6167	15
6	10007	2013-10-14	2013-10-08	parta4d5fd1	2
7	10008	2013-10-14	2013-10-08	part08cadf5	1
8	10009	2013-10-14	2013-10-08	part5cc4989	10
9	10010	2013-10-14	2013-10-08	part912ae4c	1

=====

Product line data shape: (11997, 3)

Product line data columns: ['partID', 'productLine', 'prodCategory']

First few rows of product line data:

	partID	productLine	prodCategory
0	part00005ba	line4c	Liquids
1	part000b57d	line61	Machines
2	part00123bf	linec1	Marketable
3	part0021fc9	line61	Machines
4	part0027e86	line2f	Machines
5	part002ed95	line4c	Liquids
6	part0030856	lineb8	Machines
7	part0033dfd	line49	Liquids
8	part0037a2a	linea3	Marketable
9	part003caee	linea3	Marketable

Understanding the Data

Shipments Data: Contains individual line items for each shipment, including: - **shipID:** Unique identifier for each shipment - **partID:** Product identifier - **plannedShipDate:** When the shipment was supposed to go out - **actualShipDate:** When it actually shipped - **quantity:** How many units were shipped

Product Category and Line Data: Contains product category information: - **partID:** Links to shipments data - **productLine:** The category each product belongs to - **prodCategory:** The category each product belongs to

Business Questions We'll Answer: 1. Does service level (on-time shipments) vary across product categories? 2. How often do orders include products from more than one category?

The Seven Mental Models: A Progressive Learning Journey

Now we'll work through each of the seven mental models using method chaining, starting simple and building complexity.

1. Assign: Adding New Variables

Mental Model: Create new columns with calculations and transformations.

Let's start by calculating whether each shipment was late:

```
# Simple assignment - calculate if shipment was late
shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days
    )
)

print("Added lateness calculations:")
print(shipments_with_lateness[['shipID', 'plannedShipDate', 'actualShipDate', 'is_late', 'days_late']])
```

Added lateness calculations:

	shipID	plannedShipDate	actualShipDate	is_late	days_late
0	10001	2013-11-06	2013-10-04	False	-33
1	10002	2013-10-15	2013-10-04	False	-11
2	10003	2013-10-25	2013-10-07	False	-18
3	10004	2013-10-14	2013-10-08	False	-6
4	10005	2013-10-14	2013-10-08	False	-6

Method Chaining Tip for New Python Users

Why use `lambda df`? When chaining methods, we need to reference the current state of the dataframe. The `lambda df` tells pandas “use the current dataframe in this calculation.” Without it, pandas would look for a variable called `df` that doesn't exist.

Alternative approach: You could also write this as separate steps, but method chaining keeps related operations together and makes the code more readable.

Discussion Questions: Assign Mental Model

Question 1: Data Types and Date Handling - What is the `dtype` of the `actualShipDate` series? How can you find out using code?

Answer: The `dtype` of the `actualShipDate` series is `datetime64[ns]`. I found this by running the code `print(shipments_df['actualShipDate'].dtype)`.

```
print(shipments_df['actualShipDate'].dtype)
```

datetime64[ns]

- Why is it important that both `actualShipDate` and `plannedShipDate` have the same data type for comparison? Comparison may not work if the datatypes are different. For example, if `actualShipDate` is a string and `plannedShipDate` is a datetime, the comparison will not work because the string cannot be compared to a datetime.

Question 2: String vs Date Comparison - Can you give an example where comparing two dates as strings would yield unintuitive results, e.g. what happens if you try to compare “04-11-2025” and “05-20-2024” as strings vs as dates?

Answer: Results would be unintuitive since the strings are not in the correct format. If you try to compare “04-11-2025” and “05-20-2024” as strings, the comparison will not work because the strings are not in the correct format.

```
print("04-11-2025" > "05-20-2024")
```

False

Returns False because the strings are not in the correct format.

```
print(datetime.strptime("04-11-2025", "%m-%d-%Y") > datetime.strptime("05-20-2024", "%m-%d-%Y"))
```

True

Returns True because the strings are in the correct format.

Question 3: Debug This Code

```
# This code has an error - can you spot it?
shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days,
        lateStatement="Darn Shipment is Late" if shipments_df['is_late'] else "Shipment is on time"
    )
)
```

What’s wrong with the `lateStatement` assignment and how would you fix it?

Answer: The `lateStatement` assignment is wrong as the `is_late` column is a boolean series, not a single boolean value.

```

shipments_with_lateness = (
    shipments_df
    .assign(
        is_late=lambda df: df['actualShipDate'] > df['plannedShipDate'],
        days_late=lambda df: (df['actualShipDate'] - df['plannedShipDate']).dt.days,
        lateStatement=lambda df: np.where(
            df['actualShipDate'] > df['plannedShipDate'],
            "Darn Shipment is Late",
            "Shipment is on Time"
        )
    )
)
print(shipments_with_lateness['lateStatement'])

```

```

0      Shipment is on Time
1      Shipment is on Time
2      Shipment is on Time
3      Shipment is on Time
4      Shipment is on Time
...
3995    Shipment is on Time
3996    Shipment is on Time
3997    Darn Shipment is Late
3998    Shipment is on Time
3999    Shipment is on Time
Name: lateStatement, Length: 4000, dtype: object

```

Understanding the Methods

- `.query()`: Query rows based on conditions (like SQL WHERE clause)
- `.filter()`: Filter to keep specific columns by name
- **Alternative:** You could use `.loc[]` for more complex row querying, but `.query()` is often more readable

Discussion Questions: Subset Mental Model

Question 1: Query vs Boolean Indexing - What's the difference between using `.query('is_late == True')` and `[df['is_late'] == True]`?
 - Which approach is more readable and why?

Question 2: Additional Row Querying - Can you show an example of using a variable like `late_threshold` to query rows for shipments that are at least `late_threshold` days late, e.g. what if you wanted to query rows for shipments that are at least 5 days late?

Briefly Give Answers to the Discussion Questions In This Section

Question 1: Query vs Boolean Indexing

Answer: There is no difference in results between using `.query('is_late == True')` and `[df['is_late'] == True]` is that `.query('is_late == True')` is more readable and easier to understand.

Question 2: Additional Row Querying

Answer: `late_threshold` can be used to query rows for shipments that are at least `late_threshold` days late.

```
late_threshold = 5
print(shipments_with_lateness[shipments_with_lateness['days_late'] >= late_threshold])
```

	shipID	plannedShipDate	actualShipDate	partID	quantity	is_late	\
776	10192	2013-10-09	2013-10-14	part0164a70	2	True	
777	10192	2013-10-09	2013-10-14	part9259836	1	True	
778	10192	2013-10-09	2013-10-14	part4526c73	1	True	
779	10192	2013-10-09	2013-10-14	partbb47e81	2	True	
780	10192	2013-10-09	2013-10-14	part008482f	1	True	
...	
3896	10956	2013-09-24	2013-10-15	part98c1c48	1	True	
3897	10956	2013-09-24	2013-10-15	part82e69e9	1	True	
3898	10956	2013-09-24	2013-10-15	partf23fd1e	2	True	
3899	10956	2013-09-24	2013-10-15	part825873c	1	True	
3997	11001	2013-10-04	2013-10-15	partd4952a8	1	True	

	days_late	lateStatement
776	5	Darn Shipment is Late
777	5	Darn Shipment is Late
778	5	Darn Shipment is Late
779	5	Darn Shipment is Late
780	5	Darn Shipment is Late
...
3896	21	Darn Shipment is Late
3897	21	Darn Shipment is Late
3898	21	Darn Shipment is Late
3899	21	Darn Shipment is Late
3997	11	Darn Shipment is Late

[186 rows x 8 columns]

Returns the rows where the `days_late` column is greater than or equal to `late_threshold`.

3. Drop: Removing Unwanted Data

Mental Model: Remove columns or rows you don't need.

Let's clean up our data by removing unnecessary columns:

```
# Create a cleaner dataset by dropping unnecessary columns
clean_shipments = (
    shipments_with_lateness
    .drop(columns=['quantity']) # Drop quantity column (not needed for our analysis)
    .dropna(subset=['plannedShipDate', 'actualShipDate']) # Remove rows with missing dates
)

print(f"Cleaned dataset: {len(clean_shipments)} rows, {len(clean_shipments.columns)} columns")
print("Remaining columns:", clean_shipments.columns.tolist())
```

Cleaned dataset: 4000 rows, 7 columns

Remaining columns: ['shipID', 'plannedShipDate', 'actualShipDate', 'partID', 'is_late', 'days_late']

Discussion Questions: Drop Mental Model

Question 1: Drop vs Filter Strategies - What's the difference between `.drop(columns=['quantity'])` and `.filter()` with a list of columns you want to keep? - When would you choose to drop columns vs filter to keep specific columns?

Question 2: Handling Missing Data - What happens if you use `.dropna()` without specifying subset? How is this different from `.dropna(subset=['plannedShipDate', 'actualShipDate'])`? - Why might you want to be selective about which columns to check for missing values?

Briefly Give Answers to the Discussion Questions In This Section

Question 1: Drop vs Filter Strategies

Answer: The difference between `.drop(columns=['quantity'])` and `.filter()` with a list of columns you want to keep is that `.drop(columns=['quantity'])` drops the `quantity` column from the dataframe, while `.filter()` keeps the columns in the list.

```
print(clean_shipments.columns.tolist())
```

['shipID', 'plannedShipDate', 'actualShipDate', 'partID', 'is_late', 'days_late', 'latency']

Returns the columns in the dataframe.

```
print(clean_shipments.filter(['shipID', 'partID', 'plannedShipDate', 'actualShipDate'],
```

	shipID	partID	plannedShipDate	actualShipDate	days_late
0	10001	part92b16c5	2013-11-06	2013-10-04	-33
1	10002	part66983b	2013-10-15	2013-10-04	-11
2	10003	part8e36f25	2013-10-25	2013-10-07	-18
3	10004	part30f5de0	2013-10-14	2013-10-08	-6
4	10005	part9d64d35	2013-10-14	2013-10-08	-6
...
3995	10999	part0275794	2013-10-15	2013-10-15	0
3996	10999	part04da31b	2013-10-15	2013-10-15	0
3997	11001	partd4952a8	2013-10-04	2013-10-15	11
3998	11002	parta8850c7	2013-10-15	2013-10-15	0
3999	11003	part7114b3c	2013-10-15	2013-10-15	0

[4000 rows x 5 columns]

Returns the columns in the list.

Question 2: Handling Missing Data

Answer: Not all columns having missing values need to be dropped all times. If you only need to check for missing values in certain columns. For example, if you only need to check for missing values in the `plannedShipDate` and `actualShipDate` columns, you can use `.dropna(subset=['plannedShipDate', 'actualShipDate'])`.

4. Sort: Arranging Data

Mental Model: Order data by values or indices.

Let's sort by lateness to see the worst offenders:

```
# Sort by days late (worst first)
sorted_by_lateness = (
    clean_shipments
    .sort_values('days_late', ascending=False) # Sort by days_late, highest first
    .reset_index(drop=True) # Reset index to be sequential
)

print("Shipments sorted by lateness (worst first):")
print(sorted_by_lateness[['shipID', 'partID', 'days_late', 'is_late']].head(10))
```

Shipments sorted by lateness (worst first):

	shipID	partID	days_late	is_late
0	10956	part795d1a4	21	True
1	10956	partf23fd1e	21	True
2	10956	partc653823	21	True
3	10956	partb6208b5	21	True
4	10956	parte820e31	21	True
5	10956	part50c6b9a	21	True
6	10956	part1fedfcf	21	True
7	10956	part3017fa1	21	True
8	10956	part66bb851	21	True
9	10956	partd5b19e4	21	True

Discussion Questions: Sort Mental Model

Question 1: Sorting Strategies - What's the difference between `ascending=False` and `ascending=True` in sorting? - How would you sort by multiple columns (e.g., first by `is_late`, then by `days_late`)?

Abhijeet's answer:

Question 2: Index Management - Why do we use `.reset_index(drop=True)` after sorting? - What happens to the original index when you sort? Why might this be problematic?

Briefly Give Answers to the Discussion Questions In This Section

Question 1: Sorting Strategies

Answer: `ascending=False` sorts in descending order, `ascending=True` sorts in ascending order.

```
print(sorted_by_lateness.sort_values('days_late', ascending=False))
```

	shipID	plannedShipDate	actualShipDate	partID	is_late	days_late	\
0	10956	2013-09-24	2013-10-15	part795d1a4	True	21	
12	10956	2013-09-24	2013-10-15	part825873c	True	21	
1	10956	2013-09-24	2013-10-15	partf23fd1e	True	21	
21	10956	2013-09-24	2013-10-15	part54d1a21	True	21	
20	10217	2013-09-23	2013-10-14	part2081be9	True	21	
...	
3974	10059	2013-11-05	2013-10-11	part5bd94b1	False	-	
25							
3975	10059	2013-11-05	2013-10-11	partb5c97bc	False	-	
25							

3976 25	10059	2013-11-05	2013-10-11	part0945ee9	False	-
3977 25	10059	2013-11-05	2013-10-11	part5f63c90	False	-
3999 33	10001	2013-11-06	2013-10-04	part92b16c5	False	-

```

lateStatement
0      Darn Shipment is Late
12     Darn Shipment is Late
1      Darn Shipment is Late
21     Darn Shipment is Late
20     Darn Shipment is Late
...
3974   Shipment is on Time
3975   Shipment is on Time
3976   Shipment is on Time
3977   Shipment is on Time
3999   Shipment is on Time

```

[4000 rows x 7 columns]

Returns the dataframe sorted by `days_late` in descending order.

```
print(sorted_by_lateness.sort_values('days_late', ascending=True))
```

	shipID	plannedShipDate	actualShipDate	partID	is_late	days_late	\
3999 33	10001	2013-11-06	2013-10-04	part92b16c5	False	-	
3977 25	10059	2013-11-05	2013-10-11	part5f63c90	False	-	
3976 25	10059	2013-11-05	2013-10-11	part0945ee9	False	-	
3975 25	10059	2013-11-05	2013-10-11	partb5c97bc	False	-	
3974 25	10059	2013-11-05	2013-10-11	part5bd94b1	False	-	
...	
18	10956	2013-09-24	2013-10-15	parta27d449	True	21	
20	10217	2013-09-23	2013-10-14	part2081be9	True	21	
21	10956	2013-09-24	2013-10-15	part54d1a21	True	21	
11	10956	2013-09-24	2013-10-15	part82e69e9	True	21	

```

0      10956      2013-09-24      2013-10-15  part795d1a4      True      21
      lateStatement
3999      Shipment is on Time
3977      Shipment is on Time
3976      Shipment is on Time
3975      Shipment is on Time
3974      Shipment is on Time
...
18      Darn Shipment is Late
20      Darn Shipment is Late
21      Darn Shipment is Late
11      Darn Shipment is Late
0      Darn Shipment is Late

```

[4000 rows x 7 columns]

Returns the dataframe sorted by `days_late` in ascending order.

- Use lists to sort by multiple columns

```

sorted_shipments = clean_shipments.sort_values(
    by=['is_late', 'days_late'],
    ascending=[False, False] # both in descending order
)

```

Question 2: Index Management

Answer: - use `.reset_index(drop=True)` to get a clean, continuous index after sorting.
 - Without it, the old index values remain and can make the table look messy or cause confusion when referencing rows.

```

sorted_shipments = sorted_shipments.reset_index(drop=True)
print("\nAfter resetting index:")
print(sorted_shipments)

```

After resetting index:

	shipID	plannedShipDate	actualShipDate	partID	is_late	days_late	\
0	10217	2013-09-23	2013-10-14	part2081be9	True	21	
1	10956	2013-09-24	2013-10-15	part54d1a21	True	21	
2	10956	2013-09-24	2013-10-15	part0666061	True	21	
3	10956	2013-09-24	2013-10-15	parta27d449	True	21	

4	10956	2013-09-24	2013-10-15	partc63f9bc	True	21
...
3995	10059	2013-11-05	2013-10-11	part0945ee9	False	-
25						
3996	10059	2013-11-05	2013-10-11	part5bd94b1	False	-
25						
3997	10059	2013-11-05	2013-10-11	part0d00ec6	False	-
25						
3998	10059	2013-11-05	2013-10-11	part9b3abf3	False	-
25						
3999	10001	2013-11-06	2013-10-04	part92b16c5	False	-
33						

	lateStatement
0	Darn Shipment is Late
1	Darn Shipment is Late
2	Darn Shipment is Late
3	Darn Shipment is Late
4	Darn Shipment is Late
...	...
3995	Shipment is on Time
3996	Shipment is on Time
3997	Shipment is on Time
3998	Shipment is on Time
3999	Shipment is on Time

[4000 rows x 7 columns]

5. Aggregate: Summarizing Data

Mental Model: Calculate summary statistics across groups or the entire dataset.
Let's calculate overall service level metrics:

```

# Calculate overall service level metrics
service_metrics = (
    clean_shipments
    .agg({
        'is_late': ['count', 'sum', 'mean'], # Count total, count late, calculate percentage
        'days_late': ['mean', 'max'] # Average and maximum days late
    })
    .round(3)
)

print("Overall Service Level Metrics:")
print(service_metrics)

# Calculate percentage on-time directly from the data
on_time_rate = (1 - clean_shipments['is_late'].mean()) * 100
print(f"\nOn-time delivery rate: {on_time_rate:.1f}%")

```

Overall Service Level Metrics:

	is_late	days_late
count	4000.000	NaN
sum	456.000	NaN
mean	0.114	-0.974
max	NaN	21.000

On-time delivery rate: 88.6%

Discussion Questions: Aggregate Mental Model

Question 1: Boolean Aggregation - Why does `sum()` work on boolean values? What does it count?

Briefly Give Answers to the Discussion Questions In This Section

Answer: `sum()` works on boolean values because it counts the number of True values (as 1) in the series.

```
print(clean_shipments['is_late'].sum())
```

456

Returns the number of True values in the `is_late` column.

6. Merge: Combining Information

Mental Model: Join data from multiple sources to create richer datasets.

Now let's analyze service levels by product category. First, we need to merge our data:

```
# Merge shipment data with product line data
shipments_with_category = (
    clean_shipments
    .merge(product_line_df, on='partID', how='left') # Left join to keep all shipments
    .assign(
        category_late=lambda df: df['is_late'] & df['prodCategory'].notna() # Only count a
    )
)

print("\nProduct categories available:")
print(shipments_with_category['prodCategory'].value_counts())
```

Product categories available:

```
prodCategory
Marketable      1850
Machines         846
SpareParts       767
Liquids          537
Name: count, dtype: int64
```

Discussion Questions: Merge Mental Model

Question 1: Join Types and Data Loss - Why does your professor think we should use `how='left'` in most cases? - How can you check if any shipments were lost during the merge?

Question 2: Key Column Matching - What happens if there are duplicate `partID` values in the `product_line_df`?

Briefly Give Answers to the Discussion Questions In This Section

Question 1: Join Types and Data Loss

Answer: - Using `how='left'` keeps all rows from the left DataFrame (e.g., shipments) even if there's no match in the right DataFrame (e.g., productLine). This prevents losing shipments that don't have matching product information. - Check for data loss by comparing row counts or counting NaN values after the merge.

```

print("Before merge:", len(clean_shipments))
print("After merge:", len(shipments_with_category))
missing = shipments_with_category["prodCategory"].isna().sum()
print("Shipments without matching category:", missing)

```

```

Before merge: 4000
After merge: 4000
Shipments without matching category: 0

```

Question 2: Key Column Matching

Answer: Duplicate partID values cause row multiplication — each shipment gets repeated for every matching duplicate entry. This can inflate data and affect counts or aggregations, so it's best to deduplicate before merging.

7. Split-Apply-Combine: Group Analysis

****Mental Model:**** Group data and apply functions within each group.

Now let's analyze service levels by category:

```

::: {#mental-model-7-groupby .cell execution_count=20}
``` {python .cell-code}
Analyze service levels by product category
service_by_category = (
 shipments_with_category
 .groupby('prodCategory') # Split by product category
 .agg({
 'is_late': ['any', 'count', 'sum', 'mean'], # Count, late count, percentage late
 'days_late': ['mean', 'max'] # Average and max days late
 })
 .round(3)
)

print("Service Level by Product Category:")
print(service_by_category)

```

Service Level by Product Category:

	is_late				days_late	
	any	count	sum	mean	mean	max
prodCategory						
Liquids	True	537	22	0.041	-0.950	19
Machines	True	846	152	0.180	-1.336	21

Marketable	True	1850	145	0.078	-0.804	21
SpareParts	True	767	137	0.179	-1.003	21

### ! Discussion Questions: Split-Apply-Combine Mental Model

**Question 1: GroupBy Mechanics** - What does `.groupby('prodCategory')` actually do? How does it “split” the data? - Why do we need to use `.agg()` after grouping? What happens if you don't?

**Question 2: Multi-Level Grouping** - Explore grouping by `['shipID', 'prodCategory']`? What question does this answer versus grouping by `'prodCategory'` alone? (HINT: There may be many rows with identical shipID's due to a particular order having multiple partID's.)

## Briefly Give Answers to the Discussion Questions In This Section

### Answer for Question 1: GroupBy Mechanics

- `.groupby()` splits the data into category-based groups and `.agg()` combines them again with summary statistics.
- Without `.agg()`, no calculations are performed — it just have the grouping structure.

### Answer for Question 2: Multi-Level Grouping

- Grouping by `['shipID', 'prodCategory']` breaks the data down by both shipment and category, letting to see per-shipment category details.
- while grouping by `'prodCategory'` alone only shows overall category trends.

## Answering A Business Question

**Mental Model:** Combine multiple data manipulation techniques to answer complex business questions.

Let's create a comprehensive analysis by combining shipment-level data with category information:

```
Create a comprehensive analysis dataset
comprehensive_analysis = (
 shipments_with_category
 .groupby(['shipID', 'prodCategory']) # Group by shipment and category
 .agg({
 'is_late': 'any', # True if any item in this shipment/category is late
 'days_late': 'max' # Maximum days late for this shipment/category
 })
)
```

```

 })
 .reset_index()
 .assign(
 has_multiple_categories=lambda df: df.groupby('shipID')['prodCategory'].transform('n
)
)

print("Comprehensive analysis - shipments with multiple categories:")
multi_category_shipments = comprehensive_analysis[comprehensive_analysis['has_multiple_categ
print(f"Shipments with multiple categories: {multi_category_shipments['shipID'].nunique()}")
print(f"Total unique shipments: {comprehensive_analysis['shipID'].nunique()}")
print(f"Percentage with multiple categories: {multi_category_shipments['shipID'].nunique() /

```

```

Comprehensive analysis - shipments with multiple categories:
Shipments with multiple categories: 232
Total unique shipments: 997
Percentage with multiple categories: 23.3%

```

### ! Discussion Questions: Answering A Business Question

**Question 1: Business Question Analysis** - What business question does this comprehensive analysis answer?

**Question 2: Multi-Level Grouping**

- How does grouping by ['shipID', 'prodCategory'] differ from grouping by just 'prodCategory'?
- What insights can ZappTech's management gain from knowing the percentage of multi-category shipments?

### Briefly Give Answers to the Discussion Questions In This Section

**Question 1: Business Question Analysis**

Answer: This comprehensive analysis answers the business question of how many shipments have multiple product categories.

```

print(comprehensive_analysis['has_multiple_categories'].sum())

```

594

Returns the number of shipments with multiple product categories.

### Question 2: Multi-Level Grouping

Answer: - Grouping by ['shipID', 'prodCategory'] differs from grouping by just 'prodCategory' in that it breaks the data down by both shipment and category, letting you see per-shipment category details. While grouping by 'prodCategory' alone only shows overall category trends. - The insights that ZappTech's management can gain from knowing the percentage of multi-category shipments is that they can see how many shipments have multiple product categories and what categories are most common. This can help them identify which categories are most problematic and need to be addressed.

## Student Analysis Section: Mastering Data Manipulation

**Your Task:** Demonstrate your mastery of the seven mental models through comprehensive discussion and analysis. The bulk of your grade comes from thoughtfully answering the discussion questions for each mental model. See below for more details.

### Core Challenge: Discussion Questions Analysis

**For each mental model, provide:** - Clear, concise answers to all discussion questions - Code examples where appropriate to support your explanations

#### ! Discussion Questions Requirements

**Complete all discussion question sections:** 1. **Assign Mental Model:** Data types, date handling, and debugging 2. **Subset Mental Model:** Filtering strategies and complex queries 3. **Drop Mental Model:** Data cleaning and quality management 4. **Sort Mental Model:** Data organization and business logic 5. **Aggregate Mental Model:** Summary statistics and business metrics 6. **Merge Mental Model:** Data integration and quality control 7. **Split-Apply-Combine Mental Model:** Group analysis and advanced operations 8. **Answering A Business Question:** Combining multiple data manipulation techniques to answer a business question

### Professional Visualizations (For 100% Grade)

**Your Task:** Create a professional visualization that supports your analysis and demonstrates your understanding of the data.

**Create visualizations showing:** - Service level (on-time percentage) by product category

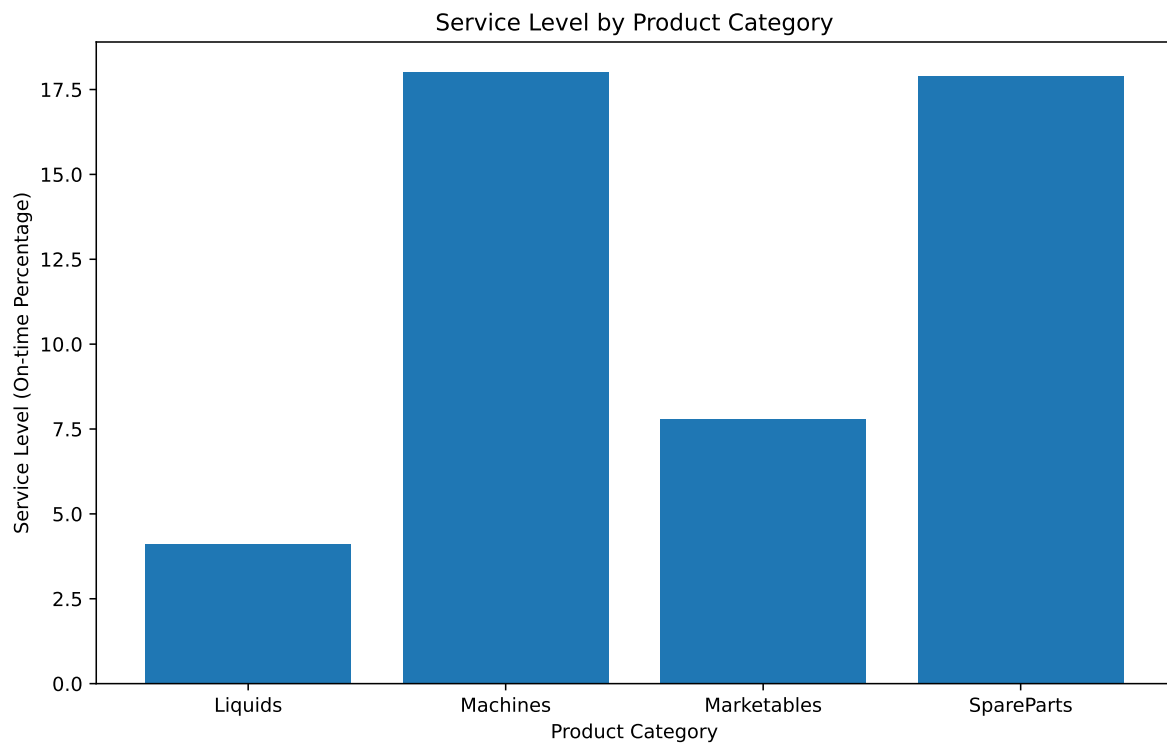
**Your visualizations should:** - Use clear labels and professional formatting - Support the insights from your discussion questions - Be appropriate for a business audience - Do not echo the code that creates the visualizations

Answer:

- Service level (on-time percentage) by product category

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 6))
plt.bar(service_by_category.index, service_by_category['is_late']['mean'] * 100)
plt.xlabel('Product Category')
plt.ylabel('Service Level (On-time Percentage)')
plt.title('Service Level by Product Category')
plt.savefig('service_level_by_product_category.png')
plt.show()
```



Returns a bar chart showing the service level (on-time percentage) by product category. This helps to identify which categories are most problematic and need to be addressed.

## Challenge Requirements

**Your Primary Task:** Answer all discussion questions for the seven mental models with thoughtful, well-reasoned responses that demonstrate your understanding of data manipulation concepts.

**Key Requirements:** - Complete discussion questions for each mental model - Demonstrate clear understanding of pandas concepts and data manipulation techniques - Write clear, business-focused analysis that explains your findings

## Getting Started: Repository Setup

### ! Getting Started

**Step 1:** Fork and clone this challenge repository - Go to the course repository and find the “dataManipulationChallenge” folder - Fork it to your GitHub account, or clone it directly - Open the cloned repository in Cursor

**Step 2:** Set up your Python environment - Follow the Python setup instructions above (use your existing venv from Tech Setup Challenge Part 2) - Make sure your virtual environment is activated and the Python interpreter is set

**Step 3:** You’re ready to start! The data loading code is already provided in this file.

**Note:** This challenge uses the same `index.qmd` file you’re reading right now - you’ll edit it to complete your analysis.

## Getting Started Tips

### i Method Chaining Philosophy

“Each operation should build naturally on the previous one”

*Think of method chaining like building with LEGO blocks - each piece connects to the next, creating something more complex and useful than the individual pieces.*

### ⚠ Important: Save Your Work Frequently!

**Before you start:** Make sure to commit your work often using the Source Control panel in Cursor (Ctrl+Shift+G or Cmd+Shift+G). This prevents the AI from overwriting your progress and ensures you don’t lose your work.

**Commit after each major step:**

- After completing each mental model section
- After adding your visualizations
- After completing your advanced method chain
- Before asking the AI for help with new code

#### **How to commit:**

1. Open Source Control panel (Ctrl+Shift+G)
2. Stage your changes (+ button)
3. Write a descriptive commit message
4. Click the checkmark to commit

*Remember: Frequent commits are your safety net!*

## **Grading Rubric**

**75% Grade:** Complete discussion questions for at least 5 of the 7 mental models with clear, thoughtful responses.

**85% Grade:** Complete discussion questions for all 7 mental models with comprehensive, well-reasoned responses.

**95% Grade:** Complete all discussion questions plus the “Answering A Business Question” section.

**100% Grade:** Complete all discussion questions plus create a professional visualization showing service level by product category.

## **Submission Checklist**

### **Minimum Requirements (Required for Any Points):**

- ☐ Created repository named “dataManipulationChallenge” in your GitHub account
- ☐ Cloned repository locally using Cursor (or VS Code)
- ☐ Completed discussion questions for at least 5 of the 7 mental models
- ☐ Document rendered to HTML successfully
- ☐ HTML files uploaded to your repository
- ☐ GitHub Pages enabled and working
- ☐ Site accessible at [https://\[your-username\].github.io/dataManipulationChallenge/](https://[your-username].github.io/dataManipulationChallenge/)

### **75% Grade Requirements:**

- ☐ Complete discussion questions for at least 5 of the 7 mental models



- ☐ Clear, thoughtful responses that demonstrate understanding
- ☐ Code examples where appropriate to support explanations

**85% Grade Requirements:**

- ☐ Complete discussion questions for all 7 mental models
- ☐ Comprehensive, well-reasoned responses showing deep understanding
- ☐ Business context for why concepts matter
- ☐ Examples of real-world applications

**95% Grade Requirements:**

- ☐ Complete discussion questions for all 7 mental models
- ☐ Complete the “Answering A Business Question” discussion questions
- ☐ Comprehensive, well-reasoned responses showing deep understanding
- ☐ Business context for why concepts matter

**100% Grade Requirements:**

- ☐ All discussion questions completed with professional quality
- ☐ Professional visualization showing service level by product category
- ☐ Professional presentation style appropriate for business audience
- ☐ Clear, engaging narrative that tells a compelling story
- ☐ Practical insights that would help ZappTech’s management

**Report Quality (Critical for Higher Grades):**

- ☐ Professional writing style (no AI-generated fluff)
- ☐ Concise analysis that gets to the point